What *Beamtime* adds

Experiment tracking for **simulations**, **training runs** and any **numerical experiment**, made for people and agents working together.

FEATURE	USUALLY	WITH BEAMTIME
Testable hypotheses *say what you expect, then check it*	Hypotheses live in your head, a doc, or buried in your coding agent's chat history.	**Proposals**, the **experiments** that test them, **notes** and **verdicts** in one place; predictions are checked for you.
People and agents, together *one shared record of every experiment*	Runs sit on a hosted server; agents reach them through a remote API.	Agents work on the run files themselves, local or remote, with **fast, compact MCP** answers. Everything in **one place**.
Notes and memory, written together *by people and agents, shown in the UI*	Notes and reports drift apart from the runs and from what was predicted.	Notes record **verdicts** next to the data; a shared **memory** keeps the big picture. All of it visible in the UI.
Views you fully control *reusable blocks; coding is not the bottleneck*	Dashboards clicked together for each project, hard to reuse or hand to an agent.	Agents build **any view** you ask for from **reusable blocks**; a good view, once made, serves every new experiment.
Plain files you own *no server, no database, no account*	Data locked in a hosted service, a database or binary event files.	Every run is a plain folder you can read, search and version. **Nothing to host**, nothing to migrate.
Built for clusters *simulations and training jobs alike*	Jobs need network to upload, or offline runs are synced by hand later.	See your cluster runs from your laptop over **ssh**: only the answers travel, and compute nodes need **no internet**.

The research cycle
what Beamtime keeps, experiment after experiment

